

## 1 IMPORTING THE NECESSARY LIBRARIES

```
[31]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## 2 UPLOADING THE DATASET(CSV FORMAT) TO THE WORKSPACE

```
[32]: data=pd.read_csv('CO2_Emissions.csv')
```

```
[33]: data
```

```
[33]:
```

	Engine Size(L)	Cylinders	Fuel Consumption City (L/100 km)	\
0	2.0	4	9.9	
1	2.4	4	11.2	
2	1.5	4	6.0	
3	3.5	6	12.7	
4	3.5	6	12.1	
...	...	...	...	
7380	2.0	4	10.7	
7381	2.0	4	11.2	
7382	2.0	4	11.7	
7383	2.0	4	11.2	
7384	2.0	4	12.2	
	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	\	
0	6.7	8.5		
1	7.7	9.6		
2	5.8	5.9		
3	9.1	11.1		
4	8.7	10.6		
...	...	...		
7380	7.7	9.4		
7381	8.3	9.9		
7382	8.6	10.3		

7383	8.3	9.9
7384	8.7	10.7

	Fuel Consumption Comb (mpg)	CO2 Emissions(g/km)
0	33	196
1	29	221
2	48	136
3	25	255
4	27	244
...	...	...
7380	30	219
7381	29	232
7382	27	240
7383	29	232
7384	26	248

[7385 rows x 7 columns]

### 3 EXPLORATORY DATA ANALYSIS

```
[34]: #I am renaming the columns as they have spaces between them, and also to improve
      ↪ readability of features
data.rename(columns={'Engine Size(L)': 'engine_size', 'Cylinders':
      ↪ 'cylinders', 'Fuel Consumption City (L/100 km)': 'fuel_cons_city',
      ↪ 'Fuel Consumption Hwy (L/100 km)': 'fuel_cons_hwy', 'Fuel
      ↪ Consumption Comb (L/100 km)': 'fuel_cons_comb',
      ↪ 'Fuel Consumption Comb (mpg)': 'fuel_cons_comb_mpg', 'CO2
      ↪ Emissions(g/km)': 'co2'
      }, inplace=True)
```

```
[35]: data.head()
```

```
[35]:   engine_size  cylinders  fuel_cons_city  fuel_cons_hwy  fuel_cons_comb  \
0         2.0         4         9.9         6.7         8.5
1         2.4         4        11.2         7.7         9.6
2         1.5         4         6.0         5.8         5.9
3         3.5         6        12.7         9.1        11.1
4         3.5         6        12.1         8.7        10.6

      fuel_cons_comb_mpg  co2
0             33  196
1             29  221
2             48  136
3             25  255
4             27  244
```

```
[36]: print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7385 entries, 0 to 7384
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   engine_size           7385 non-null   float64
1   cylinders              7385 non-null   int64
2   fuel_cons_city        7385 non-null   float64
3   fuel_cons_hwy         7385 non-null   float64
4   fuel_cons_comb        7385 non-null   float64
5   fuel_cons_comb_mpg    7385 non-null   int64
6   co2                   7385 non-null   int64
dtypes: float64(4), int64(3)
memory usage: 404.0 KB
None
```

```
[37]: data.describe()
```

```
[37]:
```

	engine_size	cylinders	fuel_cons_city	fuel_cons_hwy	\
count	7385.000000	7385.000000	7385.000000	7385.000000	
mean	3.160068	5.615030	12.556534	9.041706	
std	1.354170	1.828307	3.500274	2.224456	
min	0.900000	3.000000	4.200000	4.000000	
25%	2.000000	4.000000	10.100000	7.500000	
50%	3.000000	6.000000	12.100000	8.700000	
75%	3.700000	6.000000	14.600000	10.200000	
max	8.400000	16.000000	30.600000	20.600000	

	fuel_cons_comb	fuel_cons_comb_mpg	co2
count	7385.000000	7385.000000	7385.000000
mean	10.975071	27.481652	250.584699
std	2.892506	7.231879	58.512679
min	4.100000	11.000000	96.000000
25%	8.900000	22.000000	208.000000
50%	10.600000	27.000000	246.000000
75%	12.600000	32.000000	288.000000
max	26.100000	69.000000	522.000000


```
[38]: data.shape
```

```
[38]: (7385, 7)
```

```
[39]: data.isnull().sum()
```

```
[39]: engine_size      0
      cylinders        0
      fuel_cons_city   0
      fuel_cons_hwy    0
      fuel_cons_comb    0
      fuel_cons_comb_mpg 0
      co2              0
      dtype: int64
```

```
[40]: import seaborn as sns
      corr_matrix = data.corr()
      plt.figure(figsize=(10, 8))
      sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
      plt.title('Correlation Heatmap')
      plt.show()
```



```
[41]: # Dropping the redundant features(which are highly correlated with each other)
data.drop(columns=['fuel_cons_city', 'fuel_cons_hwy', 'fuel_cons_comb_mpg'],
           inplace=True)
```

```
[42]: data
```

```
[42]:      engine_size  cylinders  fuel_cons_comb  co2
0           2.0         4           8.5  196
1           2.4         4           9.6  221
2           1.5         4           5.9  136
3           3.5         6          11.1  255
4           3.5         6          10.6  244
...
7380         2.0         4           9.4  219
7381         2.0         4           9.9  232
7382         2.0         4          10.3  240
7383         2.0         4           9.9  232
7384         2.0         4          10.7  248
```

[7385 rows x 4 columns]

```
[43]: x=data.iloc[:, :-1].values
      y=data.iloc[:, -1].values
```

## 4 SCALING THE NECESSARY FEATURES

```
[44]: from sklearn.preprocessing import StandardScaler
      sc=StandardScaler()
      x=sc.fit_transform(x)
```

```
[45]: x.shape
```

```
[45]: (7385, 3)
```

```
[46]: x
```

```
[46]: array([[ -0.85672099, -0.88340757, -0.85574185],
        [ -0.5613172 , -0.88340757, -0.47542306],
        [ -1.22597573, -0.88340757, -1.75467716],
        ...,
        [ -0.85672099, -0.88340757, -0.23340202],
        [ -0.85672099, -0.88340757, -0.37169976],
        [ -0.85672099, -0.88340757, -0.09510428]])
```

## 5 SPLITTING THE DATASET INTO TRAINING SET AND TEST SET

```
[47]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
```

## 6 Training the Multiple Linear Regression model on the Training set

```
[48]: from sklearn.linear_model import LinearRegression
      regressor=LinearRegression()
      regressor.fit(x_train,y_train)
```

```
[48]: LinearRegression()
```

## 7 Predicting the Test set results

```
[49]: y_pred=regressor.predict(x_test)
```

```
[50]: y_pred
```

```
[50]: array([243.31052887, 200.93236303, 256.83307024, ..., 228.66076015,
          170.66557481, 220.75232578])
```

## 8 EVALUATING THE MODEL NOW

```
[51]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, \
      ↪root_mean_squared_error
      mae = mean_absolute_error(y_test, y_pred)
      mse = mean_squared_error(y_test, y_pred)
      rmse = root_mean_squared_error(y_test, y_pred)
      r2 = r2_score(y_test, y_pred)

      # Print the evaluation metrics
      print(f"Mean Absolute Error: {mae:.2f}")
      print(f"Mean Squared Error: {mse:.2f}")
      print(f"Root Mean Squared Error: {rmse:.2f}")
      print(f"R-squared: {r2:.2f}")
```

Mean Absolute Error: 13.56

Mean Squared Error: 421.59

Root Mean Squared Error: 20.53

R-squared: 0.88

## 9 VISUALISING THE RESULT

```
[52]: import matplotlib.pyplot as plt
plt.scatter(y_test, y_pred)
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs Predicted')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red')
plt.show()
```

